



AI Facilitation Group Curriculum

Six biweekly sessions from September to December, with a seventh slot in January for a topic the group picks. Each session has a facilitator plan followed by a one-page participant handout. The first five sessions are sequenced so that participants can use AI as a study aid for the qualifying exams in November.

REAL WORK ONLY.

1	2	3	4	5	6	7
Prompting and Attribution	Documents and Failure	Code Over Arithmetic	Agents and Cost	Skill Docs and Quals	Build and Critique	Group's Pick

► Scroll to read

Biweekly, Thursday and Friday

September 2026 to January 2027

6 sessions and one open slot

WORKING ASSUMPTIONS —

Edit if wrong.

- 90-minute sessions; each plan flags what to cut for 60.
- Sessions are scheduled on Thursday and Friday pairs, every other week. A session that needs more room (the build session, especially) can use both days.
- Participants have enterprise access to a frontier model with file upload, code execution, and a deep-research

SCHEDULE —

DATES	SESSION	TOPIC
Sep 17 & 18	1	Prompting as Spec-Writing, and Attribution

mode. Plans are written tool-agnostic; swap in the RAND-approved product.

- Sessions run virtually. Everyone joins from a machine they can screen-share from, with one live project open (thesis chapter, RAND task, lit review, qual reading). Breakout rooms handle pair work; a shared doc or Miro board stands in for the whiteboard. Each plan carries ready-made materials for anyone who arrives without a project.
- A shared running document ("the Playbook") captures what worked, what broke, and what the model lied about. By the last session it is the group's own RAND-specific reference.

FACILITATION NOTES —

- Assign the cold-open demo at the *end* of the prior session, not the start of the current one.
- When someone asks a question you don't know, put it to the room first, then to the model, then to the Playbook as an open item.
- Keep a "Lies Log" section in the Playbook. Every confident, wrong model output goes in it with the prompt that produced it and its bin. This becomes the best teaching material you have.

CHANGES FROM THE FIRST DRAFT —

From reviewer notes of September 10.

- Sessions reordered so documents and code come before agents, and the foundation is in place before quals.
- Attribution and disclosure (old Session 7) now sit in Session 1, since they are ground rules.
- The failure taxonomy and pre-flight checklist (old Session 8) now live in Session 2's handout and in every debrief.
- The editing passes (old Session 7) now sit in Handout 6, as a way to evaluate a build's output.
- The critic moves (old Session 5) now run inside the build session as the critique format.
- Tiering, effort, and cost (old Session 2) now pair with agents in Session 4.

DATES	SESSION	TOPIC
Oct 1 & 2	2	Documents, Context Rot, and How Models Fail
Oct 15 & 16	3	Code Over Arithmetic
Oct 29 & 30	4	Agents, Subagents, and Cost
Nov 12 & 13	5	Skill Docs, with Quals in View
Nov 26 & 27	No session	Quals and Thanksgiving
Dec 10 & 11	6	Build Something, and the Critic in the Room
Dec 24 & 25	No session	Holiday
Jan 7 & 8	7	Topic the group picks

STANDING STRUCTURE —

OPEI	CORE	DEBRI		
------	------	-------	--	--

One 90-minute session. Cold open, core, debrief, homework.

BLOCK	TIME	WHAT HAPPENS
Cold open	10 min	One pre-assigned person demos something they did with AI since last session. Real work only.

BLOCK	TIME	WHAT HAPPENS
Core	45–60 min	One technique, hands-on, on each person's own project. Facilitator drops into breakout rooms.
Debrief	10–15 min	What worked, what broke, where the model lied. Every failure gets a bin from the taxonomy in Handout 2. Someone logs it in the Playbook.
Homework	2 min	Apply the technique to something real before next session.