

1

SESSION — Sep 17 & 18

Prompting as Spec-Writing, and Attribution

10 · 55 · 10



GOAL —

Participants stop writing one-line prompts and start writing specifications: context, constraints, examples, output format, success criteria. They also leave knowing what RAND's AI-use policy says about attribution and disclosure.

WHY FIRST —

Everything downstream (documents, data, agents) fails quietly if the ask is vague. Most people who "know the basics" have never seen their own prompt history critiqued. And nobody should spend a term guessing what they are allowed to do.

READY-MADE MATERIALS —

Use these when nobody has their own. Everything here is public or was written for this course.

The policy

Screen-share the current RAND AI-use policy and read the relevant passages aloud.

Bad prompt for the cold open

Summarize this report. Run it on the GAO High-Risk Series update (free at [gao.gov/high-risk-list](https://www.gao.gov/high-risk-list); also the document for Session 2).

The rewrite

I'm a RAND policy analyst drafting a two-page brief for congressional appropriations staff. Attached is GAO's latest High-Risk Series update. Summarize what changed since the previous update: areas added, removed, or with a changed rating. Under 250 words. Only claims GAO itself makes; no recommendations of your own. Output: a three-column table (area, what changed, page cite), then one paragraph on the overall trend.

Three bad prompts for anyone with no history to excavate

Is this methodology sound? with nothing attached. Write a lit review on sanctions effectiveness. Make this better, followed by any paragraph.

Plan

90 MINUTES

10

MIN

COLD OPEN

You demo, since nobody has homework yet. Show one prompt you wrote badly and the rewrite. Set the tone: this group is about failure as much as success.

55

MIN

CORE

1 Policy and attribution 10

Screen-share the RAND AI-use policy. Read the relevant passages together. Answer: what must be disclosed, to whom, and when? What is prohibited? What data cannot go into which tools? Where is it silent? Nobody should leave guessing.

2 Excavation 10

Everyone opens their chat history and copies their three most recent real prompts into a scratch doc. No cleanup.

3 The five parts 10

Introduce the spec frame on a shared screen. A good prompt usually has: **Role/context** (who is the reader, what is the project), **Task** (one verb, one object), **Constraints** (length, sources, tone, what to exclude), **Examples** (one good, one bad if possible), **Output shape** (table, memo, bullets, JSON). Not every prompt needs all five; every prompt should have consciously skipped the ones it skips.

4 Rewrite 15

Each person rewrites one of their three prompts against the frame and runs both versions. Same model, same session type.

5 Pairs 10

Breakout rooms in pairs; each person shares their screen. Partner reads both outputs cold and says which is better and why, before seeing which prompt produced it.

10

MIN

DEBRIEF

Round the room: what changed most? Usually it's output shape or the missing "who is this for." Log two or three before/after pairs in the Playbook.

HOMEWORK

Rewrite one real prompt per day this week. Bring the one where the rewrite made the *smallest* difference; that's the interesting case.

60

MIN CUT

FOR 60 MINUTES

Cut the pairs step; do rewrites solo and debrief directly. Cut the policy block to five minutes and assign the rest as reading from the handout.

FACILITATOR
WATCH-OUTS

- Someone will say "I just iterate in conversation instead." Fine, but ask them to count turns. A spec upfront usually saves five.
 - Pushback on "examples" as too much work. Point out that one example of the output shape often beats three paragraphs of description.
 - Someone will ask whether using the model as an editor requires disclosure. Answer from the policy, not from opinion. If the policy is unclear, log it as a question to raise upward.
-