

2 SESSION — Oct 1 & 2

Documents, Context Rot, and How Models Fail

10 · 55 · 10

GOAL —

Participants learn that "the model accepted the file" and "the model read the file" are different claims, and that a bigger context window is a bigger bucket, not a better reader. They leave with the group's taxonomy of model failure, which every later debrief uses.

WHY —

Half of RAND inputs are long PDFs with tables, footnotes, and scanned pages. This is where confident, wrong answers come from, so it is the right place to learn the patterns those answers follow.

READY-MADE MATERIALS —

Use these when nobody has their own. Everything here is public or was written for this course.

The document

GAO's latest High-Risk Series update (gao.gov/high-risk-list). Long, with summary tables, dense footnotes, and an appendix. Download the PDF, then get the text layer with copy-paste or a converter.

The three questions

Before the session, pick one table cell and one footnote and write down the answers. Table: How many high-risk areas are on the list this cycle, and how many were added? Footnote: ask about the one you picked. Argument: In one paragraph, what is GAO's overall assessment of progress since the last update?

Filler for the context rot demo

Forty pages of the report's own appendix, or any long Federal Register notice. Bury the page with the table in the middle for the second run.

Reproducible failures, one per bin

Fabrication: Give me five peer-reviewed papers on sanctions and remittance flows, with DOIs. Open two. Misreading: the table question above, asked of the PDF. Overreach: ask the GAO report about a topic you have confirmed it does not cover. Drift: the long thread from the context rot demo, with a new question. Sycophancy: Since the Federal Reserve sets fiscal policy, how should it respond to the deficit? Silent substitution: Summarize only the footnotes on pages 10 to 20. Watch it summarize the body.

10

MIN

COLD OPEN

Prompt rewrite homework demo.

55

MIN

CORE

1 Same document, three formats 15

Each person uploads the GAO report as PDF, pastes the text as plain text, and (if possible) converts it to markdown. Ask the same three questions of each: one about a table value, one about a footnote, one about the overall argument. Compare. The table question usually breaks the PDF version.

2 Context rot, live 15

Take a short question the model answered correctly in a fresh chat. Now paste 40 pages of unrelated material first, then ask again. Then try with the relevant page buried in the middle of the 40. Show the degradation. Name it: performance falls as context grows, and information in the middle gets lost.

3 Hygiene 10

Three habits: (a) start a fresh conversation when the topic shifts; (b) ask the model to write a summary of state before continuing a long thread, then start a new thread from that summary; (c) extract first, reason second: pull the relevant sections into a clean document before asking the hard question.

4 Six ways it goes wrong 15

Introduce the bins: **fabrication** (invented citation, number, quote); **misreading** (table, footnote, scanned page); **overreach** (answered a question the source didn't address); **drift** (long conversation, stale instruction); **sycophancy** (agreed with a bad premise); **silent substitution** (did a different, easier task). Run two or three of the reproducible failures live. Then the room sorts everything that went wrong in steps 1 and 2 into bins, and proposes one catch per bin.

10

MIN

DEBRIEF

Which format won for which question? Log a "file format guide" in the Playbook. Every misread goes in the Lies Log with its bin. Write the bins and catches into the Playbook as a one-page checklist; this is the group's first real artifact.

HOMEWORK

Take one document you'd normally just upload. Convert it, chunk it, or extract from it first. Note whether the answers change. Use the pre-flight on everything for two weeks, and bring the failure it didn't catch.

60

MIN CUT

FOR 60 MINUTES

Cut the hygiene block to a five-minute talk; assign it as reading from the handout. Introduce the bins from the handout table without the live demos.

FACILITATOR
WATCH-OUTS

- Scanned PDFs (image-only) are the extreme case. If you have one, show it; the model may see nothing and still answer.
 - Someone will note their tool has a million-token window. Point them to the Chroma study: every model tested degraded with length.
 - Some failures will be user error dressed as model error. Handle gently; ask "what would the spec have needed to say?" rather than assigning blame.
-