

Agents, Subagents, and Cost

THE IDEA IN ONE LINE

Loops multiply both capability and error, so decide what a task is worth before you choose the tool.

WHEN AGENTIC HELPS

- The task has steps a human would also take (find, read, extract, compare).
- Each step produces something inspectable.
- You'd rather review a trace than do the steps yourself.

WHEN IT HURTS

- The task is one judgment call dressed up as a process.
- You can't see what the agent did between input and output.
- Speed is the only benefit, and you'll spend the saved time checking.

SUBAGENTS: THE INTERN TEST

Would you split this across five interns who can't talk to each other? Five independent reads (countries, years, documents): yes. One argument that needs shared understanding: no. Reading parallelizes. Writing and judgment don't. Parallel agents make implicit, conflicting decisions, and someone has to reconcile them: you. The real benefit is that each subagent gets a fresh context, so the main thread doesn't rot (Handout 2).

THREE TIERS

TIER	WHAT IT IS	RIGHT TOOL	CHECK/REQUIRE
A: Disposable	Reformatting, quick lookups, phrasing	Fastest, cheapest model	Glance
B: Consequential, checkable	Scoping a lit search, first-pass summaries, code you'll run, agentic extraction	Default model or extended thinking	Verify w/ you'll use
C: Judgment	Methods critique, framing, anything under your name	Best model, high effort	Treat as a draft you own

COST HAS THREE PARTS

1. Money and tokens (usually the smallest).
2. Your wall-clock time waiting.
3. Verification burden: how long it takes to confirm the output is right. This dominates for Tier C, and for every agentic run.

INSPECTION HABITS

- Read the trace. If your tool doesn't show one, treat the output as Tier C.
- Ask the agent to list every decision it made that wasn't in your instructions.
- Spot-check the step most likely to fail (usually extraction from a table or a scanned page).

PITFALLS

- Assuming parallel means thorough. It means fast.
- Letting the agent resolve a disagreement between sources. That's your call.
- Ten agents, one reviewer. The reviewer is the bottleneck.
- Using the frontier model for reformatting, or the cheap model for anything you'll cite.
- Believing that more compute fixes a bad spec. It amplifies it.

GO DEEPER

- Anthropic, "Building effective agents": anthropic.com/research/building-effective-agents
- Anthropic, "How we built our multi-agent research system" (the case for parallel reads): anthropic.com/engineering/multi-agent-research-system
- Cognition, "Don't Build Multi-Agents" (the case against parallel writes): cognition.com/blog/dont-build-multi-agents
- LangChain, "How and when to build multi-agent systems" (reconciling the two): langchain.com/blog/how-and-when-to-build-multi-agent-systems
- Simon Willison's blog (running notes on model capabilities and costs): simonwillison.net