

Skill Docs, with Quals in View

THE IDEA IN ONE LINE

Write down what you keep telling the model. Make it load automatically.

WHAT GOES IN A SKILL DOC

SECTION	EXAMPLE
When to use	For any RAND deliverable, memo, or brief.
Audience	Default reader: a policy generalist with 10 minutes.
Output conventions	Headers, no more than 3 levels. Page-cite every factual claim as (Author, p. X).
Always	Flag uncertainty. Distinguish what the source says from what you infer.
Never	Invent a statistic. Resolve a conflict between sources without flagging it.
Workflows	Lit review: extract claims to a table first, synthesize second.

FOR QUALS: THE STUDY PARTNER

Three rules do most of the work. Only the readings you attach, and "not in the readings" is an allowed answer. Quiz mode: one question at a time, no answer until you have tried. Every correction cites the page, and you check the page. Everything from Handouts 1 through 4 applies: it is a spec, the readings are documents, and the model will agree with a wrong answer unless told otherwise.

WHERE IT LIVES (DEPENDS ON YOUR TOOL)

- A Project or custom instruction block.
- A style file you paste or attach.

- A formal skill folder (SKILL.md plus optional scripts and references) for tools that support the open Agent Skills format.

KEEP IT SHORT

One page. The model reads long instructions the way it reads long documents: worse as they grow. If the file is getting long, split it: one skill per task type.

TEST IT

Run the same task with and without. If nothing changed, the skill isn't doing anything. If the output got verbose and caveat-heavy, you over-constrained.

SECURITY

A skill is a set of instructions the model will follow. Only use skill files you wrote yourself or got from a source you trust, and read every file in the folder before installing.

PITFALLS

- Encoding preferences you haven't tested. Every line should have earned its place.
- One giant skill for everything. Split by task.
- Forgetting to update it. Revisit after every failure story.

GO DEEPER

- Anthropic Agent Skills documentation: platform.claude.com/docs/en/agents-and-tools/agent-skills/overview
- Anthropic skills repository, with a template: github.com/anthropics/skills
- Agent Skills open specification: agentskills.io
- Anthropic, "Effective context engineering" (why short, on-demand instructions beat long always-on ones): anthropic.com/engineering/effective-context-engineering-for-ai-agents