

Build Something, and the Critic in the Room

THE IDEA IN ONE LINE

Technique is worthless until it makes something, and a build isn't finished until someone has tried to break it.

YOUR BUILD, IN FIVE PARTS

- 1. Problem:** what recurring task did this fix?
- 2. Build:** what did you make? (skill doc, script, pipeline, template)
- 3. Division of labor:** what did the model do, what did you do, what did you check?
- 4. Failure:** what went wrong along the way?
- 5. Next:** what would you do with another week?

FOUR MOVES

MOVE	PROMPT	WHAT IT EXPOSES
Steelman the opposition	Argue the opposite thesis as persuasively as you can.	Gaps in your case
Weakest claims	Name the three weakest claims and how a hostile reviewer attacks each.	Where to add evidence or cut
What would change your mind	What evidence would overturn this conclusion?	Whether your claim is falsifiable
The argument you're not making	Write the strongest version of the case I've left out.	Missing framing

MAKING CRITIQUE SPECIFIC

- Name the reviewer: "a RAND QA reviewer," "a skeptical appropriations staffer," "a labor economist."
- Forbid praise: Do not tell me what works. Only what doesn't.
- Demand location: Quote the sentence you're criticizing.

EDITING THE OUTPUT

PASS	PROMPT	WATCH FOR
Compression	Cut 30% without losing any claim.	It cuts your best sentence
Hedge audit	Flag every hedge. Which ones does the argument depend on?	Some hedges are the finding
First 90 seconds	Rewrite the opening for a reader with 90 seconds.	Loss of nuance the reader needs

Ask for a marked-up list, not a rewrite. Track what you reject.

WHAT YOU SHOULD NOW DO DIFFERENTLY

- Write a spec, not a question. Know the disclosure rules. (Session 1)
- Extract before you reason; start fresh threads; run the pre-flight. (Session 2)
- Ask for code, not arithmetic; read the code. (Session 3)
- Parallelize reads, never judgment. Tier the task before choosing the tool. (Session 4)
- Put what you keep saying into a skill doc. (Session 5)
- Run an adversarial pass before you send. (Session 6)

WHAT YOU SHOULD STILL NOT TRUST

- Any number the model computed in prose.
- Any citation you haven't opened.
- Any table read from a PDF you haven't spot-checked.
- Any conversation older than the topic.
- Any output that agrees with you too readily.

KEEPING IT GOING

- The Playbook has an owner, chosen in Session 6.
- Add to the Lies Log when the model gets you. Someone else will thank you.
- Re-run the failure sort in six months. The failure modes will have changed.

GO DEEPER

- Anthropic engineering blog: anthropic.com/engineering
- Simon Willison: simonwillison.net
- Ethan Mollick, *One Useful Thing*: oneusefulthing.org
- RAND's quality assurance standards for peer review.
- The group Playbook.